

NGINX Web Server

SOFTWARE ARCHITECTURE

HOOMAN BEHNEJAD FARD

Overview

- ▶ nginx were written in 2002.
- ▶ In 2004 it was released to the public under the two-clause BSD license.
- ▶ was meant to be a specialized tool to achieve more performance.
- ▶ Advanced event-based mechanisms.
- ▶ modular, event-driven, asynchronous, non-blocking.
- ▶ uses event notification mechanisms

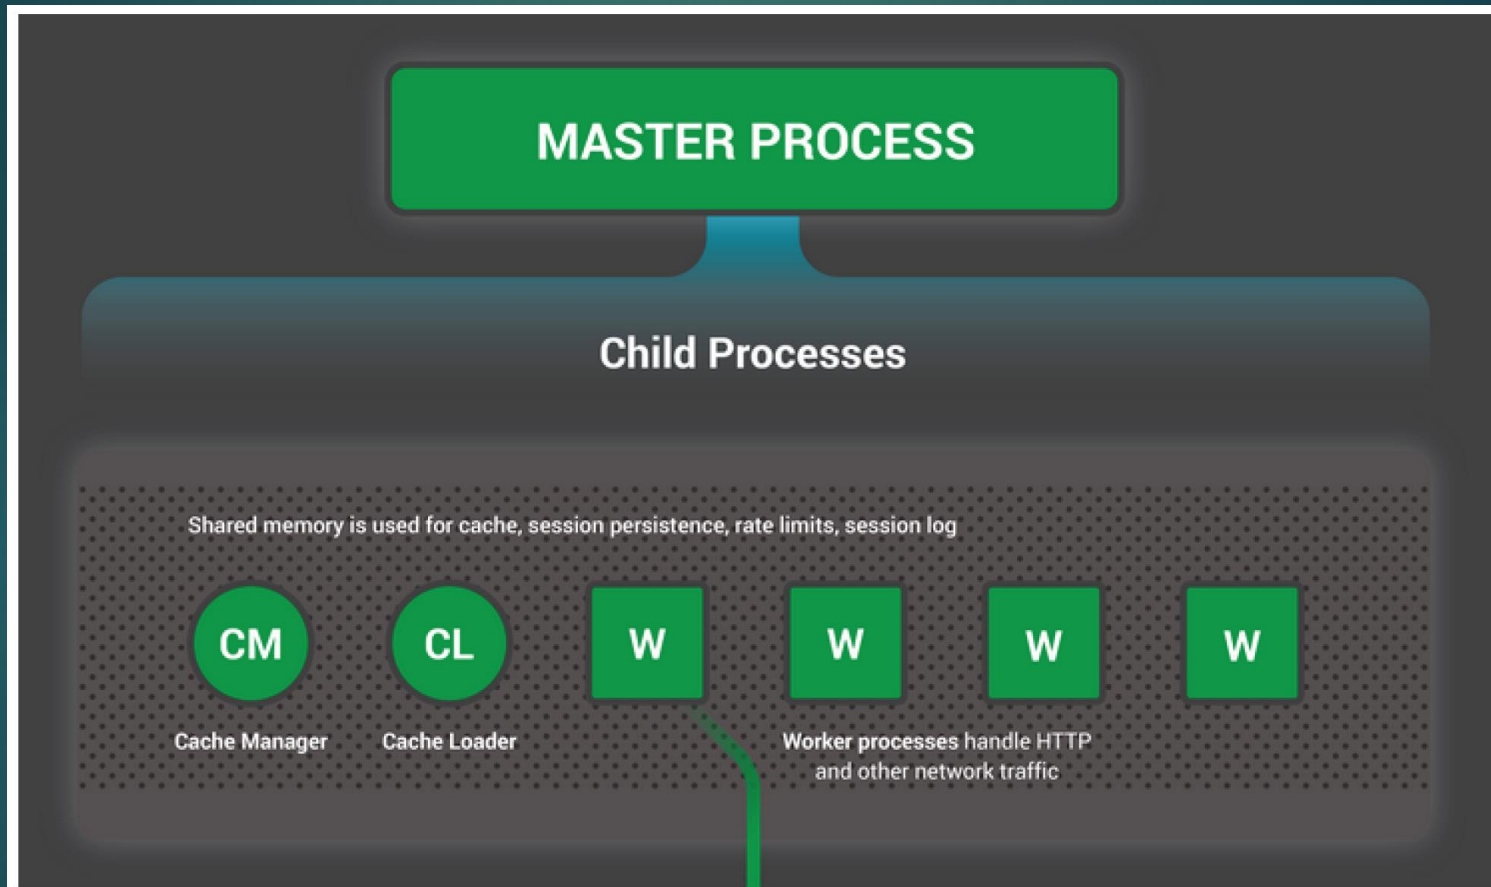
NGINX Focuses

3

- ▶ High performance
- ▶ High concurrency
- ▶ Low memory usage

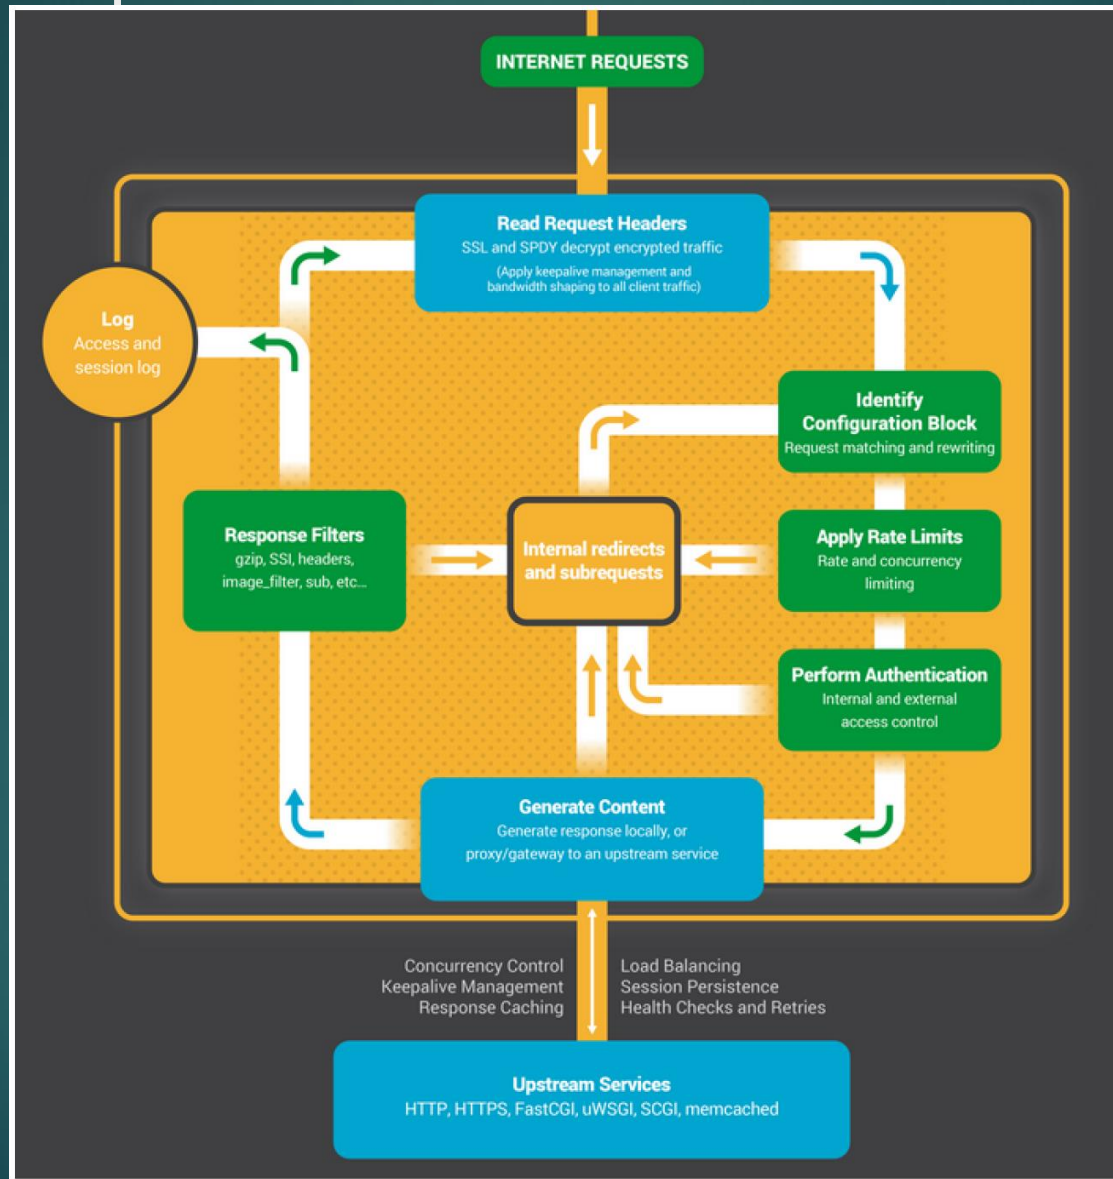
nginx's process model

4



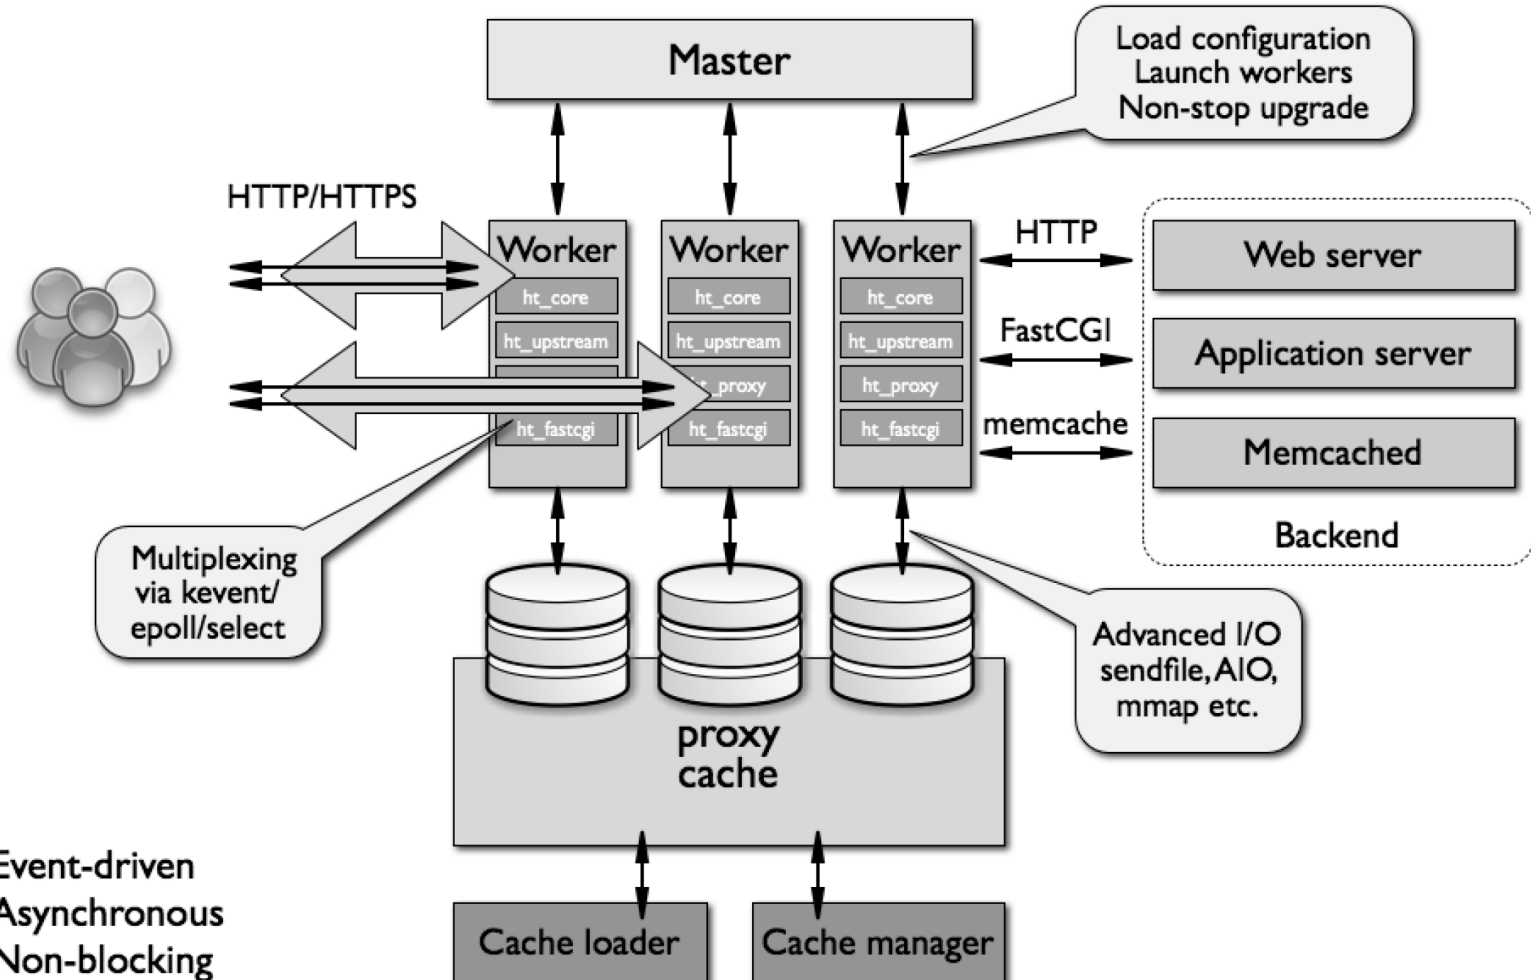
Request Flow

5



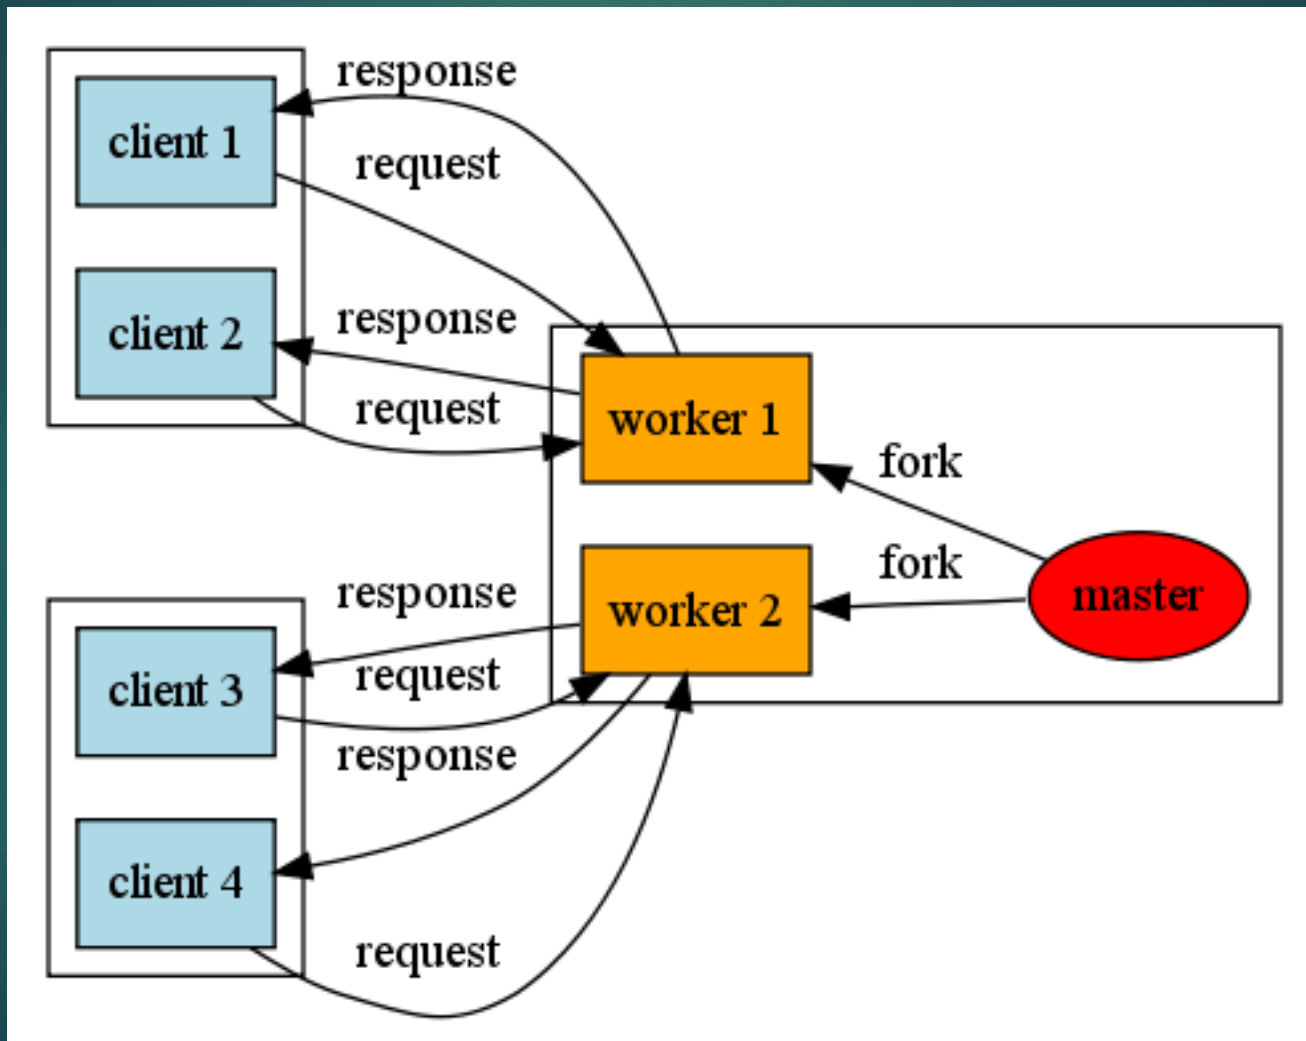
nginx's architecture

6



Master and Workers

7



NGINX uses a Non-Blocking “Event-Driven” architecture

Listen Sockets & Connection Sockets



→ Wait for an event (epoll or kqueue)

→ Event on Listen Socket:

- accept 🗯 new 📄
- set 📄 to be non-blocking
- add 📄 to the socket list

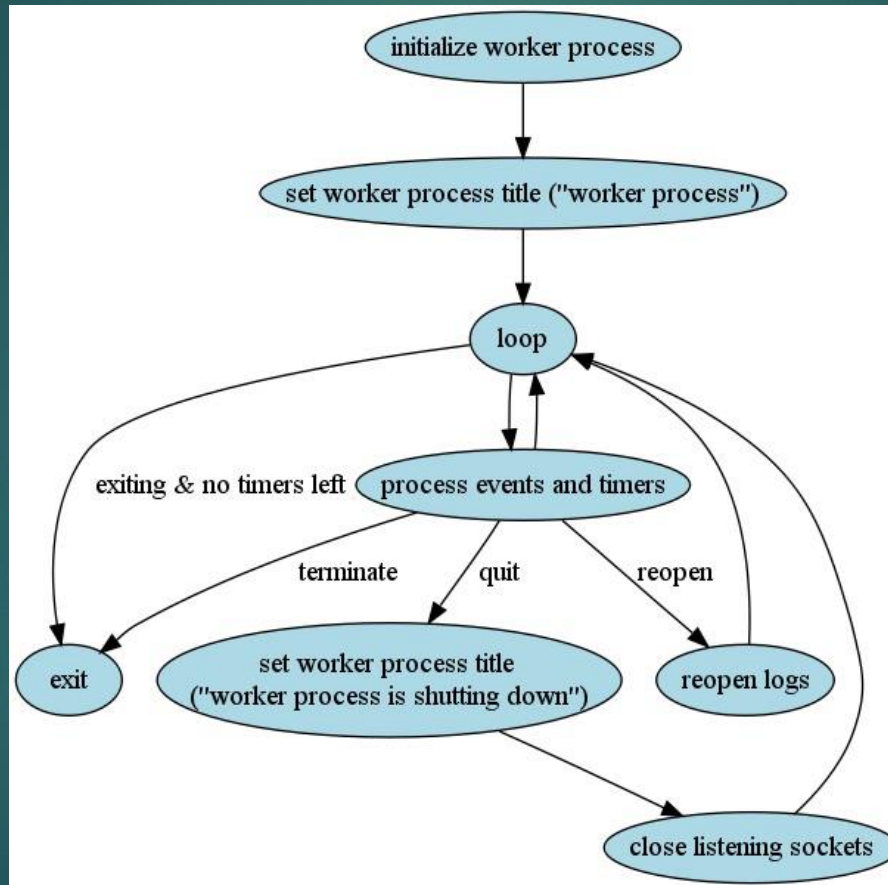
→ Event on Connection Socket:

- data in read buffer? read 📄
- space in write buffer? write 📄
- error or timeout? close 📄 & remove 📄 from socket list

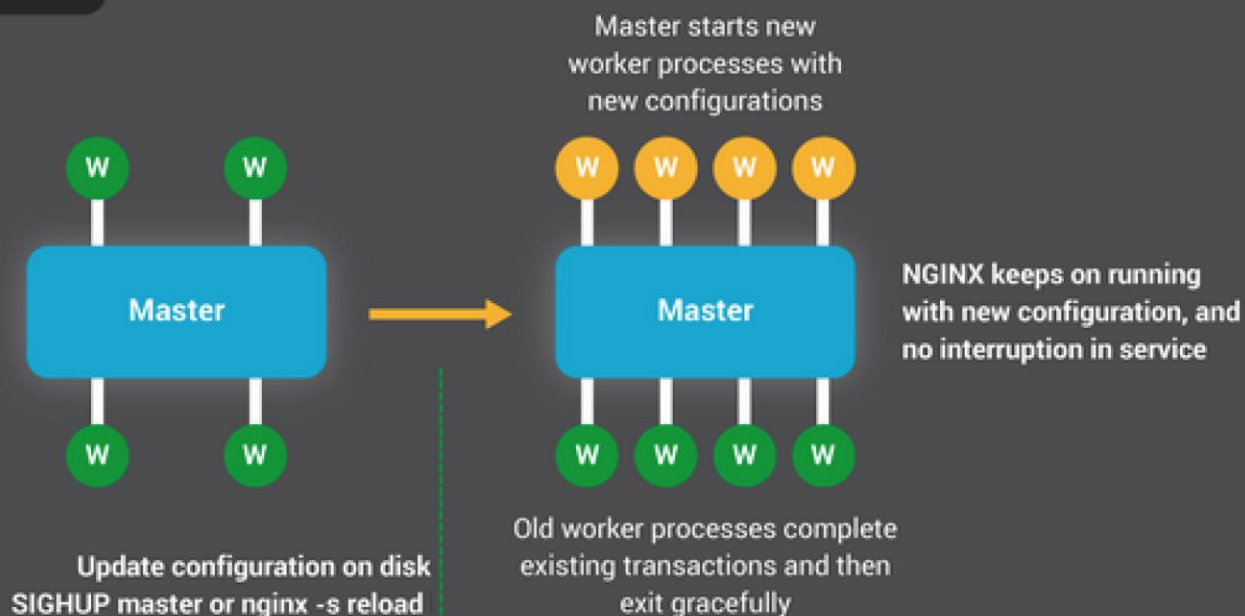
An NGINX worker can process hundreds of thousands of active connections at the same time

Worker Process Cycle

9



Load new configuration with no downtime

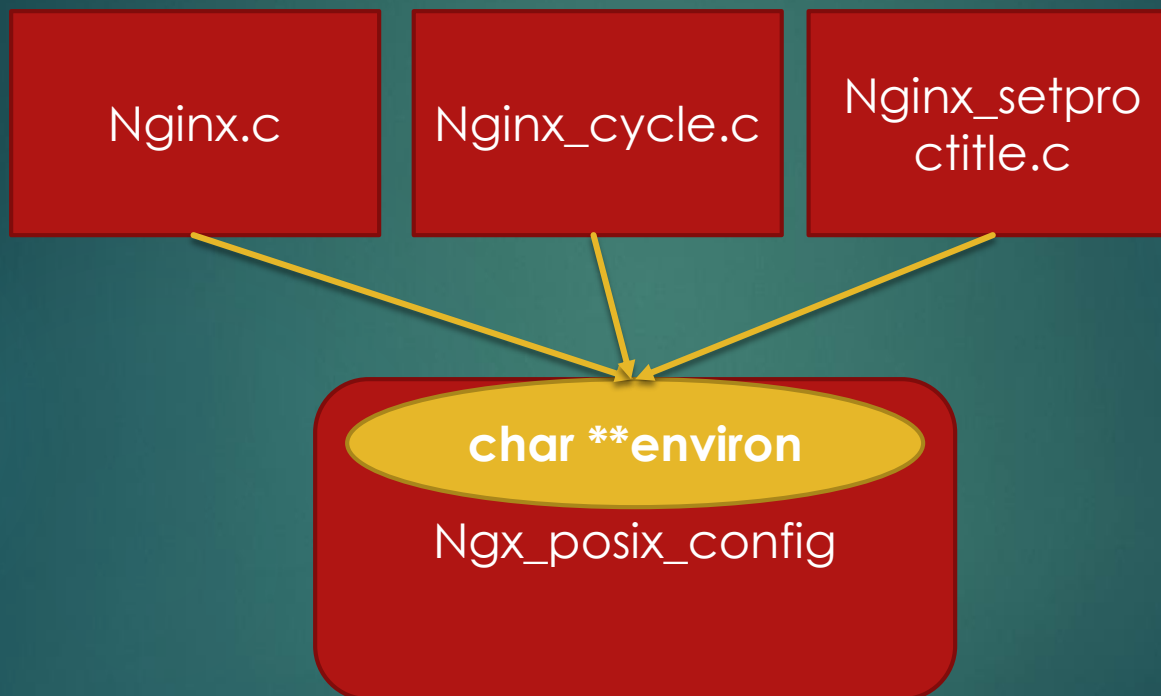


- ▶ modules are compiled along with the core at build stage.
- ▶ statically assigns each client to a specific worker thread at connection establishment.

```

- struct ngx_module_s {
    ngx_uint_t          ctx_index;
    ngx_uint_t          index;
    char                *name;
    ngx_uint_t          spare0;
    ngx_uint_t          spare1;
    ngx_uint_t          version;
    const char          *signature;
    void                *ctx;
    ngx_command_t       *commands;
    ngx_uint_t          type;
    (*init_master)(ngx_log_t *log);
    (*init_module)(ngx_cycle_t *cycle);
    (*init_process)(ngx_cycle_t *cycle);
    (*init_thread)(ngx_cycle_t *cycle);
    (*exit_thread)(ngx_cycle_t *cycle);
    (*exit_process)(ngx_cycle_t *cycle);
    (*exit_master)(ngx_cycle_t *cycle);
    uintptr_t           spare_hook0;
    uintptr_t           spare_hook1;
    uintptr_t           spare_hook2;
    uintptr_t           spare_hook3;
    uintptr_t           spare_hook4;
    uintptr_t           spare_hook5;
    uintptr_t           spare_hook6;
    uintptr_t           spare_hook7;
};

```



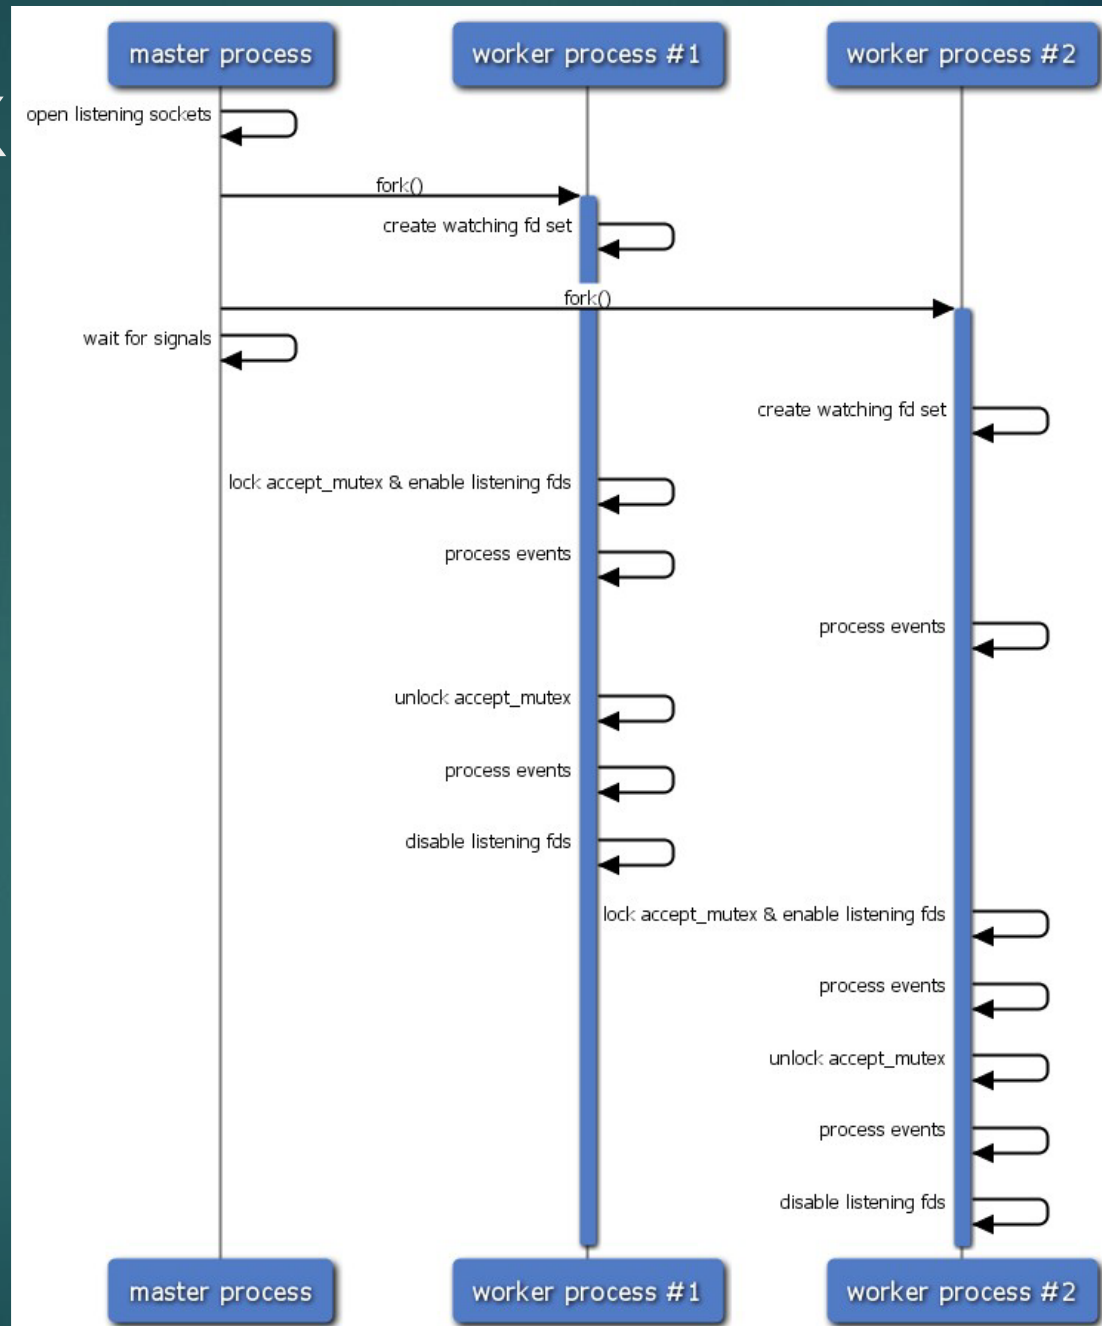
```
Char **ngx_os_environ, int ngx_kqueue, int ngx_eventfd,  
aio_context_t ngx_aio_ctx, char *ngx_os_argv_last
```

Inter-process Communication

- ▶ Signals
- ▶ Shared memory
 - ▶ Connection counter
 - ▶ stat
 - ▶ atomic & spinlock
 - ▶ Mutex

Mutex

15



Data Structures

16

- ▶ Abstract data types
 - ▶ Array
 - ▶ List
 - ▶ Queue
 - ▶ Hash table
 - ▶ Red black tree
 - ▶ Radix tree
- ▶ Characteristic
 - ▶ keep interfaces clean
 - ▶ Efficient

General Module Interface

17

- ▶ Context
 - ▶ index & ctx_index
- ▶ Directives
- ▶ Type
 - ▶ core/event/http/mail
- ▶ Hooks
 - ▶ init_master
 - ▶ called at master process initialization
 - ▶ init_module
 - ▶ called when the module is loaded
 - ▶ init_process
 - ▶ called at worker process initialization
 - ▶ exit_process
 - ▶ called at worker process termination
 - ▶ exit_master
 - ▶ called at master process termination

Event Module Interface

18

- ▶ Name
- ▶ Hooks
 - ▶ create_conf
 - ▶ init_conf
 - ▶ event_actions
 - ▶ add
 - ▶ del
 - ▶ enable
 - ▶ disable
 - ▶ add_conn
 - ▶ del_conn
 - ▶ process_changes
 - ▶ process_events
 - ▶ init
 - ▶ done

Code Organization

19

- ▶ core/
 - ▶ The backbone and infrastructure
- ▶ event/
 - ▶ The event-driven engine and modules
- ▶ http/
 - ▶ The HTTP server and modules
- ▶ mail/
 - ▶ The Mail proxy server and modules
- ▶ misc/
 - ▶ C++ compatibility test and the Google perftools module
- ▶ os/
 - ▶ OS dependent implementation files

Big Picture

20

